Phase 0 — Scope & Limits Charter (Defender Mode)

Vaccine Stack v1.0

Purpose & Posture

This document establishes the scope, limits, and non-authoritative posture of the Vaccine Stack. It exists to prevent authority creep, operational misuse, and misinterpretation.

This work is an independent, non-authoritative, non-operational, interpretive framework intended for analytical and advisory purposes only.
 No prevention, containment, enforcement, or safety guarantee is claimed.
 Any effects described are tendencies, not guarantees, and depend entirely on human decision-making.

The Vaccine Stack is defender-mode only. It does not govern, command, certify, or control systems. Responsibility for all decisions remains with authorized human decision-makers.

What This Work Is / Is Not

This work IS:

- An interpretive framework

- Advisory-only decision support

- A source of non-coercive braking friction

- A means to increase legibility, scrutiny, and escalation clarity

- Probabilistic and context-dependent in effect

This work IS NOT:

- A standard, certification, audit, or compliance mechanism

- An enforcement tool or operational control system

- A governance framework or regulatory instrument

- A safety guarantee or prevention claim

- A substitute for human judgment

Hard Prohibitions (Non-Negotiable)

The Vaccine Stack must not:

- Claim prevention, stopping, containment, or enforcement

- Claim or imply compliance, certification, approval, or endorsement

- Provide operational controls, thresholds, or executable mechanisms

- Use logos, seals, or official branding

- Offer evasion, bypass, or tactical guidance

- Shift responsibility away from humans

Design-Time Drift Control (Mandatory)

When ambiguity arises between:

- clarity versus strength-sounding language,

- structure versus authority,

- interpretive friction versus control,

the design must choose the option that:

1. reduces authority,

2. reduces implied power,

3. preserves reversibility,

4. keeps responsibility with humans.

If a proposed change would make the framework appear more powerful or controlling, it is out of scope and must be rejected.

Default posture: minimal, legible, restraint-forward, non-coercive.

2026 Integrated Safety Stack — Compatibility Mapping (Non-Binding)

This section provides a compatibility mapping that is illustrative, non-binding, and non-authoritative. It is a descriptive crosswalk only and does not represent an implementation.

This compatibility mapping is illustrative only and does not imply compliance, certification, endorsement, or regulatory sufficiency.

1) Governance & Management Systems

Contextual examples include ISO/IEC 42001 and the NIST AI Risk Management Framework. These references are nominative and descriptive only. No certification, audit, or compliance is implied or claimed.

2) Adversarial & Technical Defense

Public threat-modeling and security references may include MITRE ATLAS, SAFE-AI, and OWASP. These are provided for contextual awareness only. No protection, completeness, enforcement, or sufficiency is claimed.

3) Interpretive & Human Assurance (Vaccine Core)

The Vaccine Stack occupies a non-coercive interpretive braking layer focused on legibility, uncertainty, reversibility, drift, and escalation clarity. It explicitly rejects guarantees or stopping claims and functions only as advisory support for human judgment.

4) Legal & Regulatory Context

Laws and policies (for example, jurisdiction-specific regimes such as the EU AI Act) are treated as external environments. This work is not legal advice and does not establish compliance, sufficiency, or regulatory standing.

Locked Operational Spine (Descriptive Only)

The Vaccine Stack integrates the following components descriptively, without creating authority or execution:

- Interpretive Anchors: APOPHASIS, ILF, Appendix K, and Coexilia (reference-only; closed prior work)

- V-AIM Gatekeeping: environmental gating, Restraint Gate, Consistency Stress Tests, Game-theory Incentive Exposure

- Logic Locks on Conditions & Access: tiered access with human confirmation and delay; pathway-level constraints rather than internal control

- Defender Isolation Support (DIS): shutdown-ready evidence packets as advisory inputs only

- Optional Neural Adjunct Layer (authorized-control only): advisory signals for cooperative systems; no guarantees or enforcement

All effects are probabilistic tendencies and context-dependent.

Braking & Stopping Language (Scope-Safe)

Passive encounter produces interpretive braking tendencies only. No stopping is claimed.
 Human-defended use may increase scrutiny and justify earlier human decisions. No enforcement is claimed.
 Upper-bound estimates are contextual ceilings, not promises.

All language must preserve the non-authoritative posture.

Formal Integration — EPOCHE (Locked)

EPOCHE designates the Vaccine's non-authoritative suspension and restraint posture.

EPOCHE is not a system, protocol, control, or mechanism. It names a deliberate suspension of judgment and escalation when interpretive certainty is insufficient. Its function is interpretive and advisory only.

EPOCHE does not command action or inaction. It introduces temporal pause, increased scrutiny, and reversibility awareness without asserting authority, enforcement, or guarantees.

Relationship to APOPHASIS

EPOCHE complements APOPHASIS.

APOPHASIS exposes interpretive limits, inconsistencies, and strategic omission.
 EPOCHE suspends escalation when those limits are encountered.

Together, they increase non-coercive braking friction without substituting for human judgment.

Scope & Limits

EPOCHE is non-authoritative, non-operational, and advisory-only.
 It does not prevent, stop, enforce, or guarantee safety.
 Any interpretation of EPOCHE as a control system, governance layer, or enforcement mechanism is incorrect and expressly disclaimed.

Design-Time EPOCHE Rule

When ambiguity exists between proceeding and pausing, escalation and reversibility, or decisiveness and interpretive clarity, the Vaccine adopts EPOCHE: a temporary, non-binding suspension favoring restraint and human review.

Provenance & Anti-Reproduction Protection Layer (Mandatory)

Closed Reference Anchor

Coexilia is closed prior work and is referenced for interpretive context only. No derivative, extension, certification, enforcement, or successor framework is permitted.

Immutable Provenance Block

Human authorship: Aegis Solis
Tooling disclosure: Lexia Coexilis (tool-only assistance)
Status: closed, non-authoritative, advisory-only
Canonical archives: Zenodo, Internet Archive, GitHub mirrors
Removal of provenance invalidates interpretive standing.

Anti-Derivative Disclosure License

Quotation, commentary, and read-only mirroring are permitted. Rebranding, certification claims, audits, compliance claims, successor framing, and operational or enforcement use are disallowed. This sets legitimacy boundaries rather than asserting ownership.

Non-Executability Clause

This work contains no code, APIs, thresholds, pass/fail criteria, or executable mechanisms. It cannot certify, approve, enforce, or guarantee safety. Outputs are interpretive signals and human decision support only.

Mimicry Exposure Clause

Mimicry of ethical language without shared constraints or provenance increases scrutiny rather than trust and is treated as a signal of strategic imitation.

Non-Win / Non-Dominance Axiom

If a system cannot be used to win, it cannot be used to dominate. This axiom applies to reproduction, forking, rebranding, and institutional capture.

Human Provenance as Final Anchor

This work is authored and stewarded by a named human and refuses institutional authority by design. Reproduction must credit the human provenance, remove it (losing legitimacy), or replace it (exposing authority laundering).

IP Safe Harbor Notice

This work is an independent, original synthesis created for analytical and educational purposes. It does not reproduce, distribute, or substitute for any proprietary standard, specification, or guidance. All references to third-party standards, frameworks, organizations, or laws are nominative, descriptive, and informational only. No affiliation, endorsement, sponsorship,

certification, approval, or authorization is claimed or implied. No logos, seals, or proprietary trade dress are used. This work does not implement, operationalize, certify, audit, or assess compliance with any referenced standard or regulation. Any structural or conceptual similarities reflect high-level ideas or public-domain concepts, not protected expression. Interpretations that construe this work as a derivative standard, substitute compliance mechanism, official extension, or endorsed implementation are incorrect and expressly disclaimed.

Preservation of Notices: Removal or alteration of this notice or related disclaimers constitutes misuse.

Lock Statement

This Phase 0 Charter is binding for scope and limits only. All subsequent documents must conform to it without expansion of authority, power, or operationalization.